

Notes and Announcements

- Midterm exam: **Oct 20**, Wednesday, In Class
- Late Homeworks
 - Turn in hardcopies to Michelle.
 - DO NOT ask Michelle for extensions.
 - Note down the date and time of submission.
 - If submitting softcopy, email to 10-701 instructors list.
 - Software needs to be submitted via Blackboard.
- HW2 out today – watch email

Projects

Hands-on experience with Machine Learning Algorithms – understand when they work and fail, develop new ones!

Project Ideas online, discuss TAs, every project must have a TA mentor


- Proposal (10%): Oct 11
- Mid-term report (25%): Nov 8
- Poster presentation (20%): Dec 2, 3-6 pm, NSH Atrium
- Final Project report (45%): Dec 6

Project Proposal

- Proposal (10%): Oct 11
 - 1 pg maximum
 - Describe data set
 - Project idea (approx two paragraphs)
 - Software you will need to write.
 - 1-3 relevant papers. Read at least one before submitting your proposal.
 - Teammate. Maximum team size is 2. division of work
 - Project milestone for mid-term report? Include experimental results.

Recitation Tomorrow!

- Linear & Non-linear Regression, Nonparametric methods
- Strongly recommended!!
- Place: NSH 1507 (Note)
- Time: 5-6 pm



TK

Non-parametric methods

Kernel density estimate,
kNN classifier, kernel regression

Aarti Singh

Machine Learning 10-701/15-781
Sept 29, 2010




MACHINE LEARNING DEPARTMENT



Parametric methods

- Assume some functional form (Gaussian, Bernoulli, Multinomial, logistic, Linear) for
 - $P(X_i|Y)$ and $P(Y)$ as in Naïve Bayes
 - $P(Y|X)$ as in Logistic regression
- Estimate parameters $(\mu, \sigma^2, \theta, w, \beta)$ using MLE/MAP and plug in
- **Pro** – need few data points to learn parameters
- **Con** – Strong distributional assumptions, not satisfied in practice

Example



Hand-written digit images
projected as points on a two-dimensional (nonlinear) feature spaces

Non-Parametric methods


- Typically don't make any distributional assumptions
- As we have more data, we should be able to learn more complex models
- Let number of parameters scale with number of training data
- Today, we will see some nonparametric methods for
 - Density estimation
 - Classification
 - Regression

Histogram density estimate

Partition the feature space into distinct bins with widths Δ_i and count the number of observations, n_i , in each bin.

$$\hat{p}(x) = \frac{n_i}{n\Delta_i} \mathbf{1}_{x \in \text{Bin}_i}$$

- Often, the same width is used for all bins, $\Delta_i = \Delta$.
- Δ acts as a smoothing parameter.




Effect of histogram bin width

$$\hat{p}(x) = \frac{n_i}{n\Delta} \mathbf{1}_{x \in \text{Bin}_i}$$

$$\# \text{ bins} = 1/\Delta$$

$$\hat{p}(x) = \frac{1}{\Delta} \frac{\sum_{j=1}^n \mathbf{1}_{X_j \in \text{Bin}_x}}{n}$$



Bias of histogram density estimate:

$$\mathbb{E}[\hat{p}(x)] = \frac{1}{\Delta} P(X \in \text{Bin}_x) = \frac{1}{\Delta} \int_{z \in \text{Bin}_x} p(z) dz \approx \frac{p(x)\Delta}{\Delta} = p(x)$$

**Assuming density is roughly constant in each bin
(holds true if Δ is small)**

Bias – Variance tradeoff

- Choice of #bins


$$\# \text{ bins} = 1/\Delta$$

$\mathbb{E}[\hat{p}(x)] \approx p(x)$ if Δ is small (p(x) approx constant per bin)

$\mathbb{E}[\hat{p}(x)] \approx \hat{p}(x)$ if Δ is large (more data per bin, stable estimate)

- Bias** – how close is the mean of estimate to the truth
- Variance** – how much does the estimate vary around mean

Small Δ , large #bins  “Small bias, Large variance”

Large Δ , small #bins  “Large bias, Small variance”

Bias-Variance tradeoff

Choice of #bins

$$\hat{p}(x) = \frac{n_i}{n\Delta} \mathbf{1}_{x \in \text{Bin}_i}$$

bins = $1/\Delta$




Image src: Bishop book




Image src: Larry book

Histogram as MLE

- Class of density estimates – constants on each bin
Parameters p_j - density in bin j


Note $\sum_j p_j = 1/\Delta$ since $\int p(x)dx = 1$

- Maximize likelihood of data under probability model with parameters p_j
- Show that histogram density estimate is MLE under this model – HW/Recitation

Kernel density estimate


- Histogram – blocky estimate

$$\hat{p}(x) = \frac{1}{\Delta} \frac{\sum_{j=1}^n \mathbf{1}_{X_j \in \text{Bin}_x}}{n}$$



- Kernel density estimate aka “Parzen/moving window method”

$$\hat{p}(x) = \frac{1}{\Delta} \frac{\sum_{j=1}^n \mathbf{1}_{||X_j - x|| \leq \Delta}}{n}$$





Kernel density estimate

- $$\hat{p}(x) = \frac{1}{\Delta} \frac{\sum_{j=1}^n K\left(\frac{X_j - x}{\Delta}\right)}{n} \quad \text{more generally}$$



boxcar kernel :

$$K(x) = \frac{1}{2}I(x),$$




Gaussian kernel :

$$K(x) = \frac{1}{\sqrt{2\pi}} e^{-x^2/2}$$



Kernel density estimation

- Place small "bumps" at each data point, determined by the kernel function.
- The estimator consists of a (normalized) "sum of bumps".




Gaussian bumps (red) around six data points and their sum (blue)

- Note that where the points are denser the density estimate will have higher values.

Kernels

boxcar kernel :

$$K(x) = \frac{1}{2}I(x),$$




Any kernel function that satisfies

$$K(x) \geq 0,$$

$$\int K(x)dx = 1$$

Gaussian kernel :


$$K(x) = \frac{1}{\sqrt{2\pi}}e^{-x^2/2}$$



Kernels

boxcar kernel :

$$K(x) = \frac{1}{2}I(x),$$




Finite support

- only need local points to compute estimate

Gaussian kernel :


$$K(x) = \frac{1}{\sqrt{2\pi}}e^{-x^2/2}$$




Infinite support

- need all points to compute estimate
- But quite popular since smoother (10-702)

Choice of kernel bandwidth




True Density



Too small


Undersmoothed

Image Source:
Larry's book – All
of Nonparametric
Statistics



Just right

Just Right





Too large

Oversmoothed

**Bart-Simpson
Density**

Histograms vs. Kernel density estimation



$\Delta = h$ acts as a smoother.

Bias-variance tradeoff

- Simulations

k-NN (Nearest Neighbor) density estimation

- Histogram $\hat{p}(x) = \frac{n_i}{n\Delta} \mathbf{1}_{x \in \text{Bin}_i}$
- Kernel density est $\hat{p}(x) = \frac{n_x}{n\Delta}$


Fix Δ , estimate number of points within Δ of x (n_i or n_x) from data

Fix $n_x = k$, estimate Δ from data (volume of ball around x that contains k training pts)

- k-NN density est $\hat{p}(x) = \frac{k}{n\Delta_{k,x}}$

k-NN density estimation

$$\hat{p}(x) = \frac{k}{n\Delta_{k,x}}$$



k acts as a smoother.


Not very popular for density estimation - expensive to compute, bad estimates

But a related version for classification quite popular

...


From
Density estimation
to
Classification

k-NN classifier




k-NN classifier

Test document



k-NN classifier (k=4)



What should we predict? ... Average? Majority? Why?

k-NN classifier

- Optimal Classifier: $f^*(x) = \arg \max_y P(y|x)$
 $= \arg \max_y p(x|y)P(y)$
- k-NN Classifier: $\hat{f}_{kNN}(x) = \arg \max_y \hat{p}_{kNN}(x|y)\hat{P}(y)$
 $= \arg \max_y k_y$ **(Majority vote)**

$$\hat{p}_{kNN}(x|y) = \frac{k_y}{n_y \Delta_{k,x}}$$


k_y \longrightarrow # training pts of class y that lie within Δ_k ball

n_y \longrightarrow # total training pts of class y


$$\sum_y k_y = k$$

$$\hat{P}(y) = \frac{n_y}{n}$$


1-Nearest Neighbor (kNN) classifier




2-Nearest Neighbor (kNN) classifier



3-Nearest Neighbor (kNN) classifier



5-Nearest Neighbor (kNN) classifier



What is the best K?

Bias-variance tradeoff


Larger K \Rightarrow predicted label is more stable

Smaller K \Rightarrow predicted label is more accurate


Similar to density estimation

Choice of K - in next class ...


1-NN classifier – decision boundary



Voronoi
Diagram



k-NN classifier – decision boundary



- K acts as a smoother (Bias-variance tradeoff)
- Guarantee: For $n \rightarrow \infty$, the error rate of the 1-nearest-neighbour classifier is never more than twice the optimal error.

Case Study:

kNN for Web Classification

- Dataset
 - 20 News Groups (20 classes)
 - Download :(<http://people.csail.mit.edu/jrennie/20Newsgroups/>)
 - 61,118 words, 18,774 documents
 - Class labels descriptions

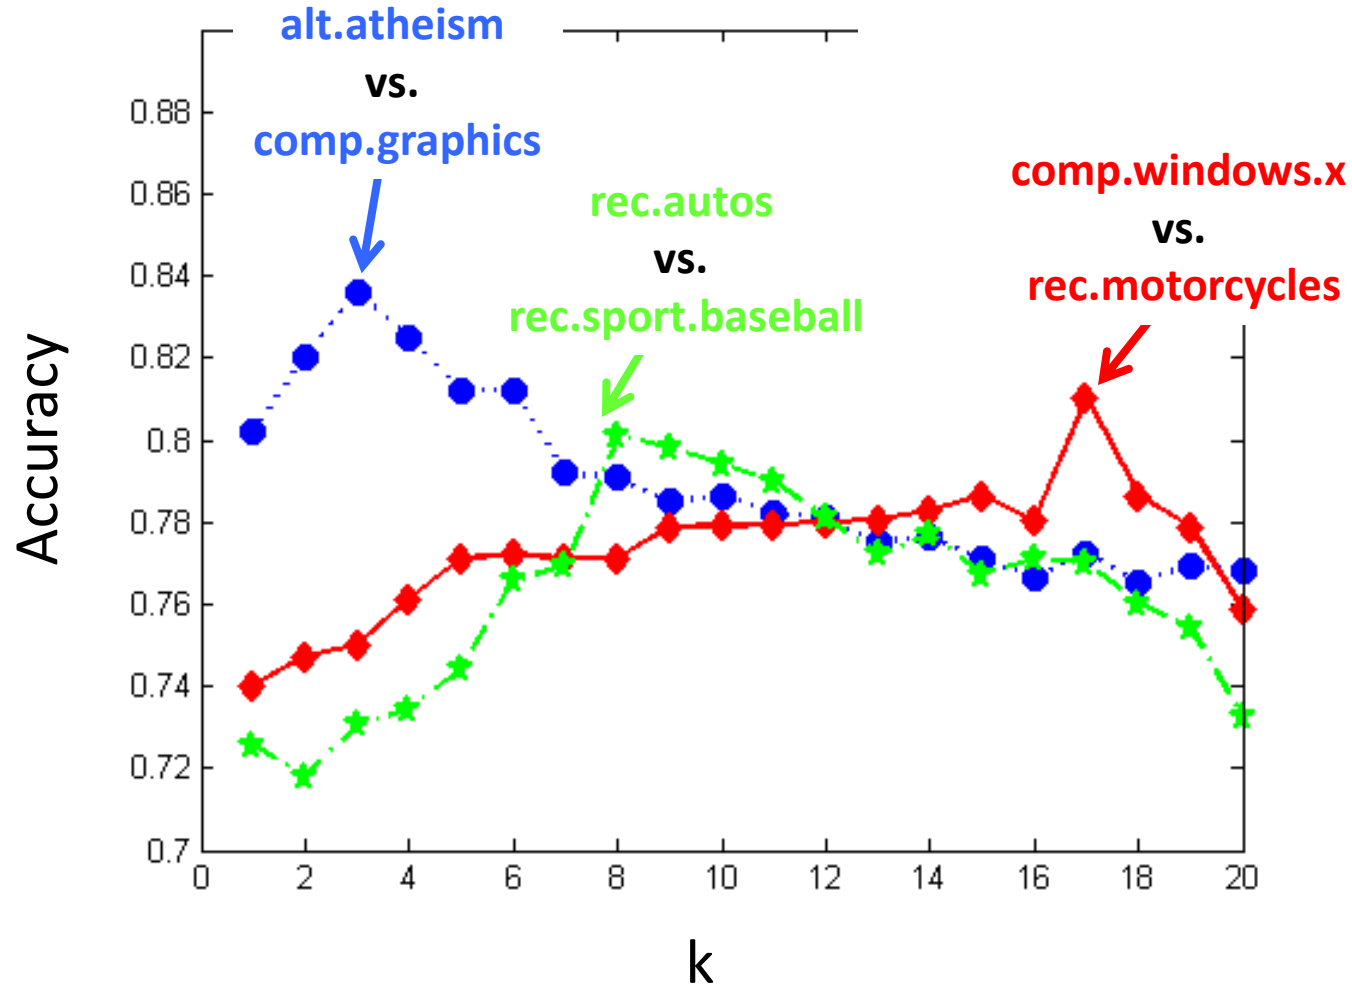
comp.graphics comp.os.ms-windows.misc comp.sys.ibm.pc.hardware comp.sys.mac.hardware comp.windows.x	rec.autos rec.motorcycles rec.sport.baseball rec.sport.hockey	sci.crypt sci.electronics sci.med sci.space
misc.forsale	talk.politics.misc talk.politics.guns talk.politics.mideast	talk.religion.misc alt.atheism soc.religion.christian

Experimental Setup

- Training/Test Sets:
 - 50%-50% randomly split.
 - 10 runs
 - report average results
- Evaluation Criteria:

$$Accuracy = \frac{\sum_{i \in \text{test set}} I(\text{predict}_i = \text{true label}_i)}{\# \text{ of test samples}}$$


Results: Binary Classes



From Classification to Regression

Temperature sensing


- What is the temperature in the room?



$$\hat{T} = \frac{1}{n} \sum_{i=1}^n Y_i$$

Average


at location x ?



$$\hat{T}(x) = \frac{\sum_{i=1}^n Y_i \mathbf{1}_{\|X_i - x\| \leq h}}{\sum_{i=1}^n \mathbf{1}_{\|X_i - x\| \leq h}}$$

"Local" Average

Kernel Regression




- Aka Local Regression
- Nadaraya-Watson Kernel Estimator

$$\hat{f}_n(X) = \sum_{i=1}^n w_i Y_i \quad \text{Where} \quad w_i(X) = \frac{K\left(\frac{X - X_i}{h}\right)}{\sum_{i=1}^n K\left(\frac{X - X_i}{h}\right)}$$

- Weight each training point based on distance to test point
- Boxcar kernel yields local average

boxcar kernel :


$$K(x) = \frac{1}{2}I(x),$$




Kernels

boxcar kernel :

$$K(x) = \frac{1}{2}I(x),$$





$$K\left(\frac{X_j - x}{\Delta}\right)$$



Gaussian kernel :

$$K(x) = \frac{1}{\sqrt{2\pi}}e^{-x^2/2}$$



Choice of kernel bandwidth h







Image Source:
Larry's book – All
of Nonparametric
Statistics

Choice of kernel is
not that important

Kernel Regression as Weighted Least Squares

$$\min_f \sum_{i=1}^n w_i (f(X_i) - Y_i)^2 \qquad w_i(X) = \frac{K\left(\frac{X - X_i}{h}\right)}{\sum_{i=1}^n K\left(\frac{X - X_i}{h}\right)}$$

Weighted Least Squares

Kernel regression corresponds to locally constant estimator obtained from (locally) weighted least squares

i.e. set $f(X_i) = \beta$ (a constant)

Kernel Regression as Weighted Least Squares

set $f(X_i) = \beta$ (a constant)

$$\min_{\beta} \sum_{i=1}^n w_i (\underbrace{\beta}_{\text{constant}} - Y_i)^2$$

$$w_i(X) = \frac{K\left(\frac{X - X_i}{h}\right)}{\sum_{i=1}^n K\left(\frac{X - X_i}{h}\right)}$$

$$\frac{\partial J(\beta)}{\partial \beta} = 2 \sum_{i=1}^n w_i (\beta - Y_i) = 0$$

Notice that $\sum_{i=1}^n w_i = 1$

$$\Rightarrow \hat{f}_n(X) = \hat{\beta} = \sum_{i=1}^n w_i Y_i$$

Local Linear/Polynomial Regression

$$\min_f \sum_{i=1}^n w_i (f(X_i) - Y_i)^2 \qquad w_i(X) = \frac{K\left(\frac{X - X_i}{h}\right)}{\sum_{i=1}^n K\left(\frac{X - X_i}{h}\right)}$$

Weighted Least Squares

Local Polynomial regression corresponds to locally polynomial estimator obtained from (locally) weighted least squares

$$\text{i.e. set } f(X_i) = \beta_0 + \beta_1(X_i - X) + \frac{\beta_2}{2!}(X_i - X)^2 + \dots + \frac{\beta_p}{p!}(X_i - X)^p$$

(local polynomial of degree p around X)

More in HW, 10-702 (statistical machine learning)

Summary

- Instance based/non-parametric approaches

Four things make a memory based learner:

1. *A distance metric, $\text{dist}(x, X_i)$*
Euclidean (and many more)
2. *How many nearby neighbors/radius to look at?*
 $k, \Delta/h$
3. *A weighting function (optional)*
 W based on kernel K
4. *How to fit with the local points?*
Average, Majority vote, Weighted average, Poly fit

Summary

- Parametric vs Nonparametric approaches

- Nonparametric models place very mild assumptions on the data distribution and provide good models for complex data

Parametric models rely on very strong (simplistic) distributional assumptions

- Nonparametric models (not histograms) requires storing and computing with the entire data set.

Parametric models, once fitted, are much more efficient in terms of storage and computation.

What you should know...

- Histograms, Kernel density estimation
 - Effect of bin width/ kernel bandwidth
 - Bias-variance tradeoff
- K-NN classifier
 - Nonlinear decision boundaries
- Kernel (local) regression
 - Interpretation as weighted least squares
 - Local constant/linear/polynomial regression