

# Decision Trees

Aarti Singh

Machine Learning 10-701/15-781  
Oct 6 , 2010



**MACHINE LEARNING** DEPARTMENT



# Learning a good prediction rule

- Learn a mapping  $f : \mathcal{X} \rightarrow \mathcal{Y}$
- Best prediction rule  $f^*(X) = \arg \min_f R(f)$
- Hypothesis space/Function class  $\mathcal{F}$ 
  - Parametric classes (Gaussian, binomial etc.)
  - Conditionally independent class densities (Naïve Bayes)
  - Linear decision boundary (Logistic regression)
  - Nonparametric class (Histograms, nearest neighbor, kernel estimators, **Decision Trees - Today**)
- Given training data, find a hypothesis/function in  $\mathcal{F}$  that is close to the best prediction rule.

$$\hat{f}_n(X) = \arg \min_{\underline{f \in \mathcal{F}}} \hat{R}_n(f) + C(f)$$

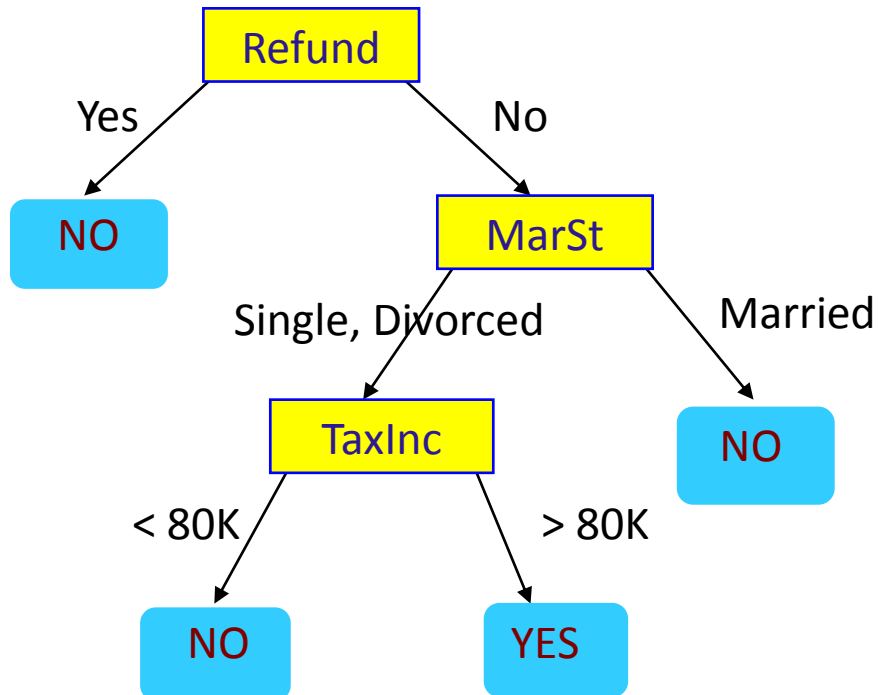
# First ...

- What does a decision tree represent
- Given a decision tree, how do we assign label to a test point

# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$



Query Data


$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

- Each internal node: test one feature  $X_i$
- Each branch from a node: selects one value for  $X_i$
- Each leaf node: predict  $Y$

# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$



Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

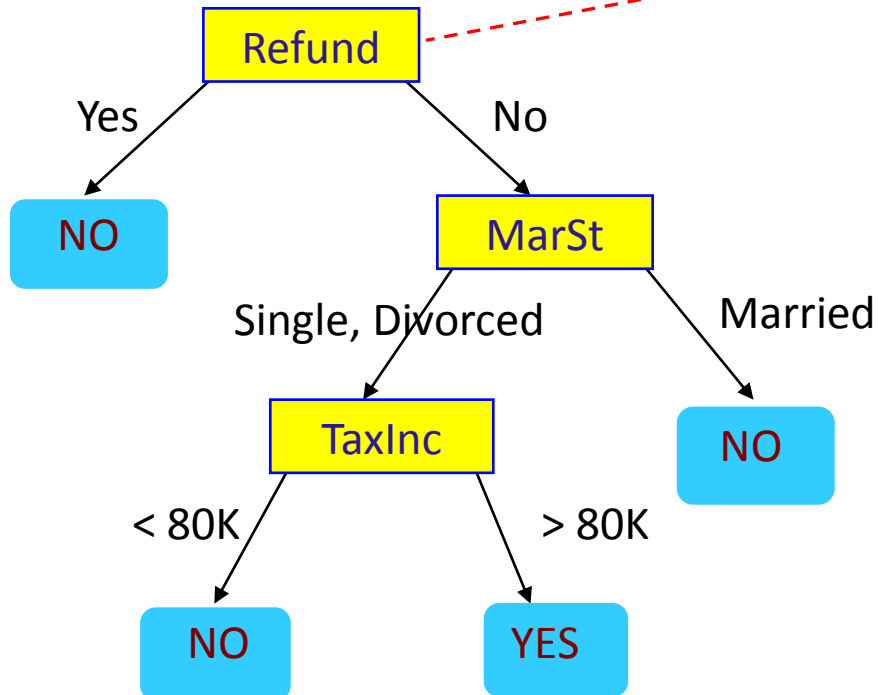
# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$

Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



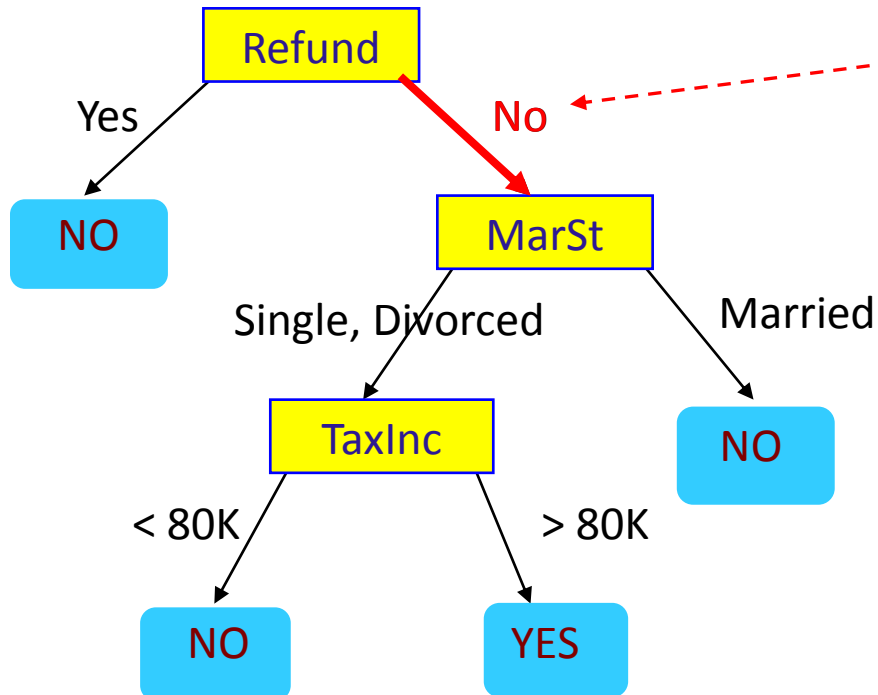
# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$

Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



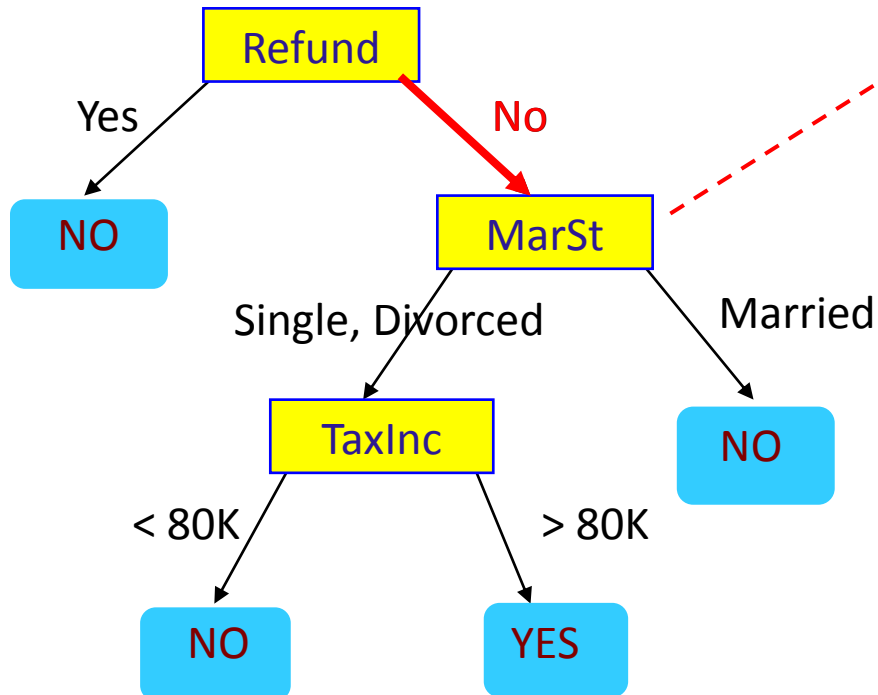
# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$

Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?





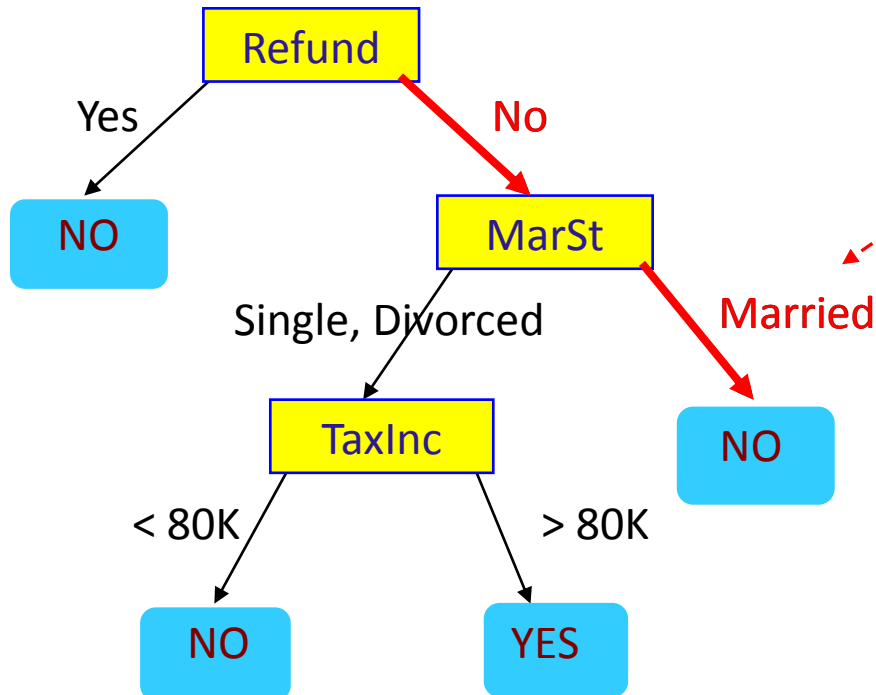
# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$

Query Data

$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



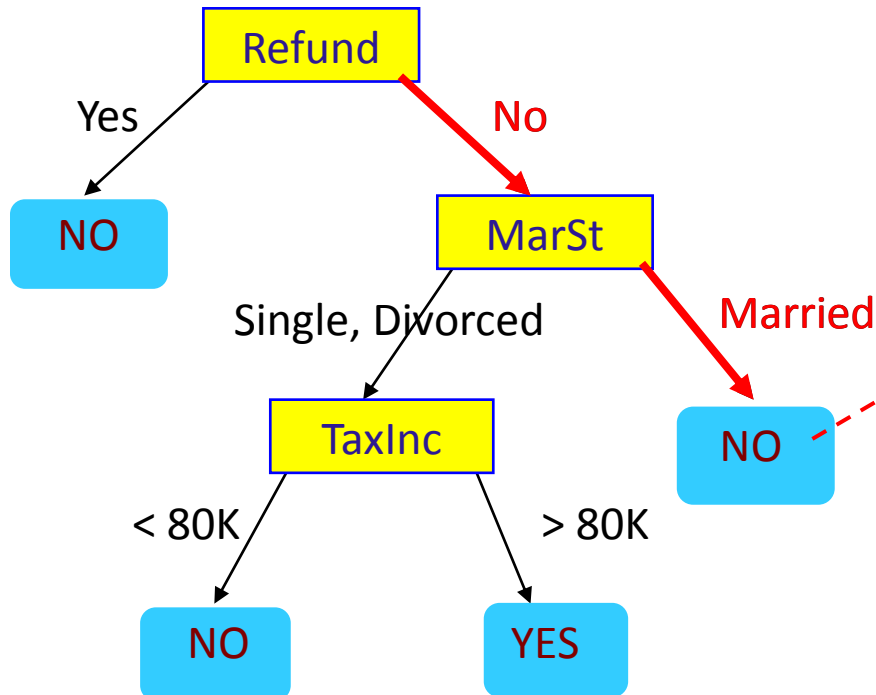
# Decision Tree for Tax Fraud Detection

$\mathcal{F}$  – Decision Trees

$$f(X_1, X_2, X_3) \in \mathcal{F}$$

Query Data



$X_1$	$X_2$	$X_3$	$Y$
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Assign Cheat to "No"

# Decision Tree more generally...

$$X_1 \geq 0.5, X_2 = \{a, b\} \text{ or } \{c, d\}$$



- Features can be discrete, continuous or categorical
- Each internal node: test some set of features  $\{X_i\}$
- Each branch from a node: selects a set of value for  $\{X_i\}$
- Each leaf node: predict  $Y$

## So far...

- What does a decision tree represent
- Given a decision tree, how do we assign label to a test point

## Now ...

- How do we learn a decision tree from training data
- What is the decision on each leaf

## So far...

- What does a decision tree represent
- Given a decision tree, how do we assign label to a test point

## Now ...

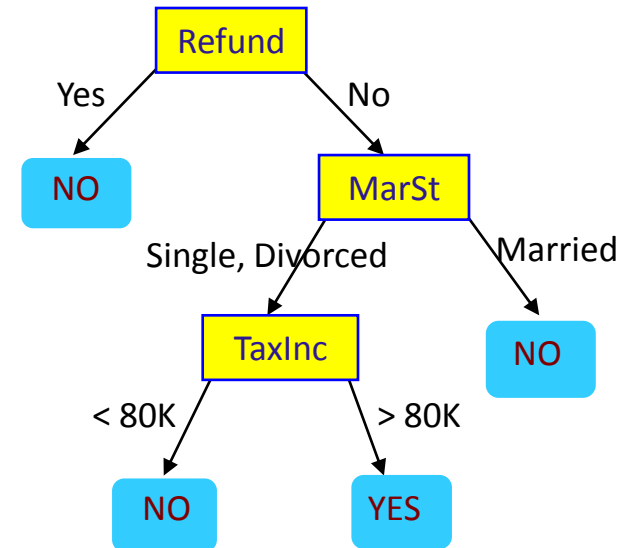
- How do we learn a decision tree from training data
- What is the decision on each leaf

# How to learn a decision tree

- Top-down induction [ID3, C4.5, CART, ...]


Main loop:

1.  $X \leftarrow$  the “best” decision attribute for next *node*
2. Assign  $X$  as decision attribute for *node*
3. For each value of  $X$ , create new descendant of *node*
4. Sort training examples to leaf nodes
5. If training examples perfectly classified, Then STOP, Else iterate over new leaf nodes





# Which feature is best to split?

$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F



Good split if we are more certain about classification after split – Uniform distribution of labels is bad

# Which feature is best to split?



Pick the attribute/feature which yields maximum information gain:

$$\arg \max_i I(Y, X_i) = \arg \max_i [H(Y) - H(Y|X_i)]$$

$H(Y)$  – entropy of  $Y$       $H(Y|X_i)$  – conditional entropy of  $Y$




# Entropy

- Entropy of a random variable  $Y$

$$H(Y) = - \sum_y P(Y = y) \log_2 P(Y = y)$$

***More uncertainty,  
more entropy!***

$Y \sim \text{Bernoulli}(p)$



**Information Theory interpretation:**  $H(Y)$  is the expected number of bits needed to encode a randomly drawn value of  $Y$  (under most efficient code)

# Andrew Moore's Entropy in a Nutshell



Low Entropy

..the values (locations of soup) sampled entirely from within the soup bowl



High Entropy

..the values (locations of soup) unpredictable... almost uniformly sampled throughout our dining room

# Information Gain

- Advantage of attribute = decrease in uncertainty

- Entropy of Y before split

$$H(Y) = - \sum_y P(Y = y) \log_2 P(Y = y)$$

- Entropy of Y after splitting based on  $X_i$

- Weight by probability of following each branch

$$\begin{aligned} H(Y | X_i) &= - \sum_x P(X_i = x) H(Y | X_i = x) \\ &= - \sum_x P(X_i = x) \sum_y P(Y = y | X_i = x) \log_2 P(Y = y | X_i = x) \end{aligned}$$

- Information gain is difference



$$I(Y, X_i) = H(Y) - H(Y | X_i)$$

**Max Information gain = min conditional entropy**

# Information Gain

$$H(Y | X_i) = - \sum_x P(X_i = x) \sum_y P(Y = y | X_i = x) \log_2 P(Y = y | X_i = x)$$

$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F



$$\hat{H}(Y|X_1) = -\frac{1}{2}[1 \log_2 1 + 0 \log_2 0] - \frac{1}{2}[\frac{1}{4} \log_2 \frac{1}{4} + \frac{3}{4} \log_2 \frac{3}{4}]$$

$$\hat{H}(Y|X_2) = -\frac{1}{2}[\frac{3}{4} \log_2 \frac{3}{4} + \frac{1}{4} \log_2 \frac{1}{4}] - \underbrace{\frac{1}{2}[\frac{1}{2} \log_2 \frac{1}{2} + \frac{1}{2} \log_2 \frac{1}{2}]}_{> 0}$$

$$\hat{H}(Y|X_1) < \hat{H}(Y|X_2)$$

# Which feature is best to split?

Pick the attribute/feature which yields maximum information gain:


$$\arg \max_i I(Y, X_i) = \arg \max_i [H(Y) - H(Y|X_i)]$$

$H(Y)$  – entropy of  $Y$        $H(Y|X_i)$  – conditional entropy of  $Y$

Feature which yields maximum reduction in entropy  
provides maximum information about  $Y$

# Expressiveness of Decision Trees



- Decision trees can express any function of the input features.
- E.g., for Boolean functions, truth table row  $\rightarrow$  path to leaf:



- There is a decision tree which perfectly classifies a training set with one path to leaf for each example
- But it won't generalize well to new examples - prefer to find more **compact** decision trees


# Decision Trees - Overfitting

One training example per leaf – overfits, need compact/pruned decision tree




# Bias-Variance Tradeoff

Ideal classifier




coarse partition




average  
classifier

bias large




Classifiers based on  
different training data


variance small




fine partition



bias small



variance large





# When to Stop?

- Many strategies for picking simpler trees:

- Pre-pruning


- Fixed depth
    - Fixed number of leaves

- Post-pruning

- Chi-square test

- Convert decision tree to a set of rules
      - Eliminate variable values in rules which are independent of label (using chi-square test for independence)
      - Simplify rule set by eliminating unnecessary rules

- Information Criteria: MDL(Minimum Description Length)



# Information Criteria

- Penalize complex models by introducing cost

$$\hat{f} = \arg \min_T \left\{ \underbrace{\frac{1}{n} \sum_{i=1}^n \text{loss}(\hat{f}_T(X_i), Y_i)}_{\text{log likelihood}} + \underbrace{\text{pen}(T)}_{\text{cost}} \right\}$$

$$\begin{aligned} \text{loss}(\hat{f}_T(X_i), Y_i) &= (\hat{f}_T(X_i) - Y_i)^2 && \text{regression} \\ &= \mathbf{1}_{\hat{f}_T(X_i) \neq Y_i} && \text{classification} \end{aligned}$$

$$\text{pen}(T) \propto |T| \quad \text{penalize trees with more leaves}$$

# Information Criteria - MDL

Penalize complex models based on their **information content**.

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \left\{ \hat{R}_n(f) + C(f) \right\}$$

$\swarrow$  # bits needed to describe  $f$   
(description length)

**MDL (Minimum Description Length)**

Example: Binary Decision trees  $\mathcal{F}_k^T = \{\text{tree classifiers with } k \text{ leafs}\}$


prefix encode each element  $f$  of  $\mathcal{F}_k^T$

$$C(f) = 3k - 1 \text{ bits}$$

$k$  leaves  $\Rightarrow 2k - 1$  nodes

$2k - 1$  bits to encode tree structure

+  $k$  bits to encode label of each leaf (0/1)



5 leaves  $\Rightarrow$  9 bits to encode structure

## So far...

- What does a decision tree represent
- Given a decision tree, how do we assign label to a test point


## Now ...

- How do we learn a decision tree from training data
- What is the decision on each leaf

# How to assign label to each leaf


Classification – Majority vote

Regression – ?




# How to assign label to each leaf

Classification – Majority vote



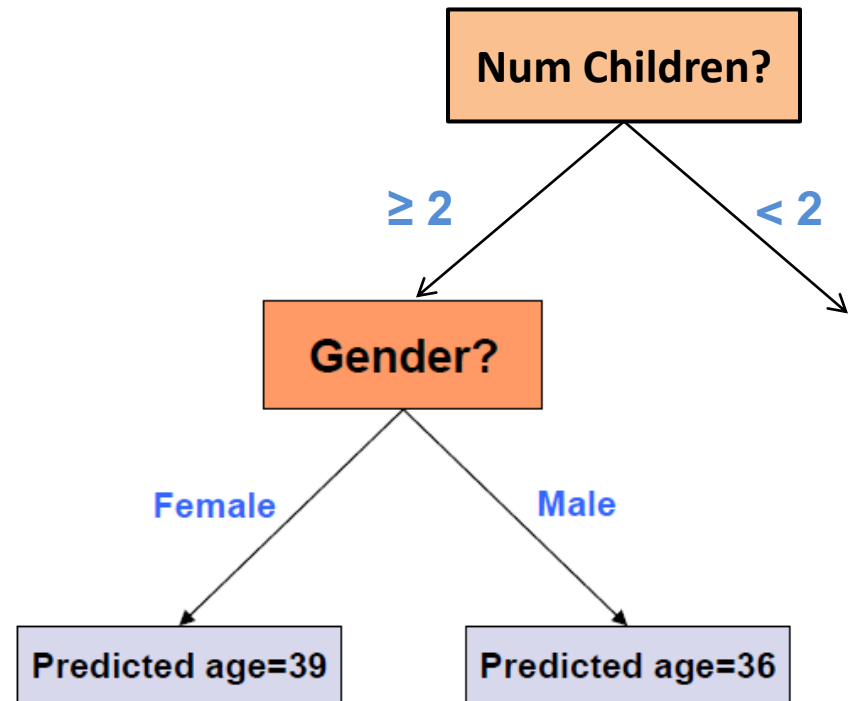
Regression – Constant/  
Linear/Poly fit



# Regression trees

$X^{(1)}$       ....       $X^{(p)}$        $Y$

Gender	Rich?	Num. Children	# travel per yr.	Age
F	No	2	5	38
M	No	0	2	25
M	Yes	1	0	72
:	:	:	:	:




Average (fit a constant ) using training data at the leaves

**Connection between  
nearest neighbor/histogram  
classifiers  
and  
decision trees**



# Local prediction


Histogram, kernel density estimation, k-nearest neighbor classifier, kernel regression



Histogram Classifier

# Local Adaptive prediction

Let neighborhood size adapt to data – small neighborhoods near decision boundary (small bias), large neighborhoods elsewhere (small variance)



Decision Tree Classifier

# Histogram Classifier vs Decision Trees

Ideal classifier





Decision tree




histogram




256 cells in  
each partition



# Application to Image Coding



1024 cells in  
each partition



# Application to Image Coding



JPEG 0.125 bpp  
non-adaptive partitioning



JPEG 2000 0.125 bpp  
adaptive partitioning

# What you should know

- Decision trees are one of the most popular data mining tools
  - Simplicity of design
  - Interpretability
  - Ease of implementation
  - Good performance in practice (for small dimensions)
- Information gain to select attributes (ID3, C4.5,...)
- Can be used for classification, regression and density estimation too
- Decision trees will overfit!!!
  - Must use tricks to find “simple trees”, e.g.,
    - Pre-Pruning: Fixed depth/Fixed number of leaves
    - Post-Pruning: Chi-square test of independence
    - Complexity Penalized/MDL model selection