

Recitation

By Field Cady
February 4, 2010

Statistics Review

Statistics vs. Probability

- Probability : deriving properties of data from the distribution
- Statistics : deducing properties of the distribution from the data
 - In this sense, Machine Learning is a form of statistics
 - But “statistics” usually refers to classical statistics

Classical Statistics Terminology

- Statistic : any function of your data
 - Can have logs, roots, arbitrary manipulation
 - Sample mean is just a function where you add and divide
- Estimator : a statistic that is intended to estimate some distribution parameter
 - Consistent estimator : enough data brings you arbitrarily close to the correct number
 - Unbiased estimator : expected value of the estimator is the actual parameter

Estimators can be biased!


- Know X is Uniform(0, θ), θ unknown
- Data = x_1, x_2, \dots, x_n
- Want to estimate θ
- Obvious estimator : $\max(x_i)$
 - Enough data points makes it close to θ
 - But it's ALWAYS an underestimate
 - You can work out the math : $\max^*(n+1)/n$ is unbiased and consistent

Subtle Example

- Estimating a normal distribution: μ and σ

- $\mu^* = (1/n)\sum x_i$

- $(\sigma^*)^2 = (1/n)\sum (x_i - \mu^*)^2$



- How good are these estimators?

- Sample mean = consistent, unbiased estimator for μ
- Sample variance
 - Consistent – converges to actual variance
 - But it is biased!
- Imagine only 2 data points.
 - μ^* exactly between them, used for calculating σ^*
 - But using any other μ^* would give higher σ^*
 - The points won't be exactly the same distance from μ
 - So we probably underestimate σ !

MLE vs Bayes Summary

MLE vs. Bayes

	MLE	Bayesian
Basic Idea	Pick distribution to make data likely	Start with belief, and adjust for data
Examples	Basic, EM alg, logistic regression	Hierarchical bayes models
Good	Data-based, often easy	Expert knowledge, mirrors thinking
Bad	Overfitting	Personal bias
Fixes	Cross-validation	Max-entropy prior

Training, Validation, and **Testing Data**

What the heck?

Reasons we do it

- Recall : this is a special case of MLE
 - MLE is subject to overfitting
- The more we train a model, the better it becomes at predicting the *training* data
- But our goal is to predict other data that is i.i.d as training data
- Idea : use some training data for testing

Cross-validation

- So we set aside some of the data as “validation data” to gauge how much to train
- But what if we randomly picked a bad validation set?
- Answer : break data into k sets, and use each set for validation separately

Testing Data

- Problem : k-fold validation isn't quite independent of validation data
 - Classifier implicitly trained on validation set, since validation data influences training cutoff
 - So validation data underestimates error
- Answer : have another chunk of data that plays no part at all in training classifier

Neural Networks

Motivation 1

- Recall linear regression :
$$\text{unit output} = \frac{1}{1 + \exp(w_0 + \sum_i w_i x_i)}$$
- Decision boundary is a straight line
- They work decently, but we don't want just linear boundaries
- Crazy idea for a hack : let's wire a bunch of them together

Motivation 2

- Human brains are fabulous
- They're made of neurons, which have multiple inputs and a single output
- Let have an embarrassingly simple input->output function and model a brain
- Note : there's work going on at using more physically realistic neural nets to get better performance

Training Neural Nets

- MLE

- Assume Gaussian errors
- Minimize sum of squares
- Backpropagation algorithm

$$W \leftarrow \arg \min_W \sum_l (y^l - \hat{f}(x^l))^2$$

- MAP :

- Add penalty term $\ln P(W)$
- If using Gaussian prior :

$$W \leftarrow \arg \min_W \left[c \sum_i w_i^2 \right] + \left[\sum_l (y^l - \hat{f}(x^l))^2 \right]$$

Regression

Several Interpretations

- Fitting a straight line to data
- A part of logistic regression
- Predicting a continuous variable
 - Rather than discrete classification
 - Will never be “exactly” right, so error rate=1
 - Instead minimize average squared error

Regularization

- General approach to problems where you penalize “extreme” solutions
- Examples
 - Regression : penalize huge parameters
 - Image denoising : penalize jitteriness
- Can be done in a principled way, but also just as a heuristic
- Can turn an under-determined set of equations into an optimization problem

What Penalty Function

- Ridge Regression : squared
- Lasso Regression : absolute value
- Important point of contention : what should penalty terms look like?

Aside : Is Least Squares Good?

- “Penalty” terms are ubiquitous
 - Least squares is just minimizing a penalty
 - Regularization penalty
 - Soon : k-means clustering minimizes squares
- Why do we use squares in the term?
 - Answer 1 : it’s easy
 - Answer 2 : it has a pretty theoretical interpretation if you use a Gaussian prior
- But is this really good?

Least Squares

- Effect of squares : emphasize outliers
 - Point 10x farther away is 100x more important
- But shouldn't we fit our model to the “normal” data points?
- Gaussians have “thin tails”
 - Few outliers, so big deviations very important

Not Squares

- But real data often has “heavy tails”
 - Major outliers more common
- Exponentials, Laplacians, etc. have more outliers
- Minimizing sum of absolute values = Laplacian prior
- Absolute value penalty function, like in Lasso, makes outliers less important
- “Robust Statistics” deals with this

Overfitting : Bias and **Variance**

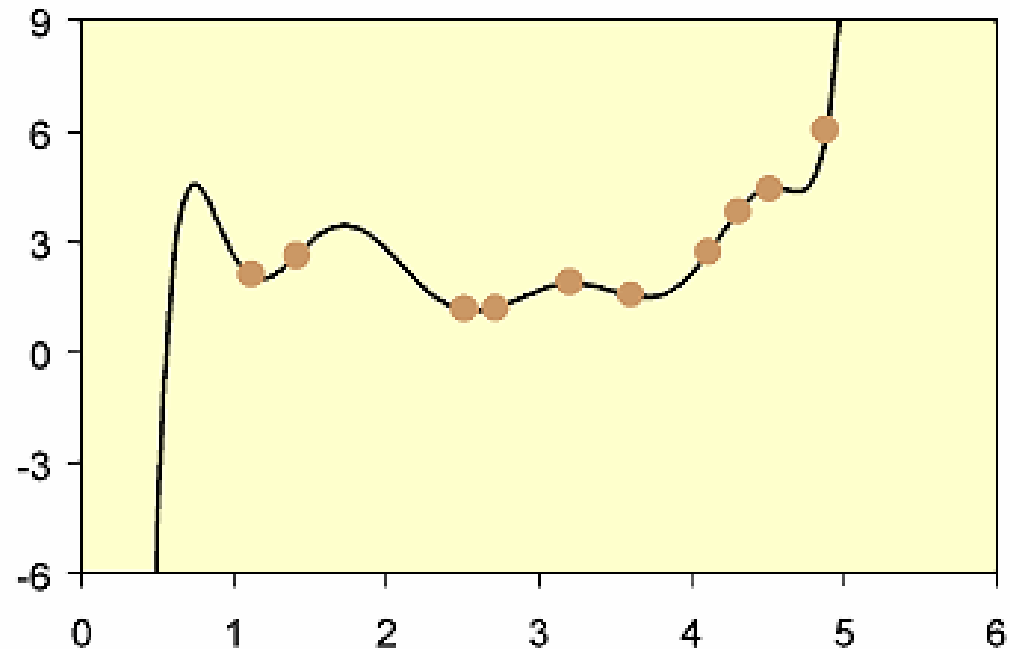
Risk and Error

- Risk for classification $P(f(X) \neq Y)$
 - Probability data is misclassified
- Error for regression $\mathbb{E}[(f(X) - Y)^2]$
 - We'll never guess continuous values exactly, but a good regression function will be close
 - (Is this a good metric? Food for thought...)

True vs. Empirical

- Empirical risk/error : how badly we do on our training data
- True risk/error : how badly we'll do on new data
- These are not the same!!!!

Example




- Empirical error : 0
- True error : very high

The Basic Conflict

- Training picks a single classifier from a “family” of classifiers
- Observation 1 : A more expressive family can represent more general distributions
- But : A more expressive family is more likely to contain a bad distribution that fits that data really well


Example

- Fitting a straight line is imperfect, but higher order polynomials are not always better...



Effect of Model Complexity

If we allow very complicated predictors, we could overfit the training data.



Empirical risk is no longer a good indicator of true risk

Excess Risk

- Some risk is inherent to the distribution
 - How close are we to this lower bound?

$$E \left[R(\hat{f}_n) \right] - R^*$$

- Lower bound independent of our data
- For this discussion, we have exactly n points

- Recall that \hat{f}_n is just a statistic of our data; it is a random variable
- Distribution depends on n
- Then its performance will be too

$$R(\hat{f}_n) = P_{XY}(\hat{f}_n(X) \neq Y)$$

$$R(\hat{f}_n) = \mathbb{E}_{XY}[(\hat{f}_n(X) - Y)^2]$$

} \hat{f}_n depends on random training dataset

Behavior of True Risk


Want predictor based on training data \hat{f}_n to be as good as optimal predictor

$$\text{Excess Risk } E[R(\hat{f}_n)] - R^* = \underbrace{\left(E[R(\hat{f}_n)] - \inf_{f \in \mathcal{F}} R(f) \right)}_{\text{estimation error}} + \underbrace{\left(\inf_{f \in \mathcal{F}} R(f) - R^* \right)}_{\text{approximation error}}$$

finite sample size
+ noise


Due to randomness
of training data

Due to restriction
of model class



Behavior of True Risk

$$E[R(\hat{f}_n)] - R^* = \underbrace{\left(E[R(\hat{f}_n)] - \inf_{f \in \mathcal{F}} R(f)\right)}_{\text{estimation error}} + \underbrace{\left(\inf_{f \in \mathcal{F}} R(f) - R^*\right)}_{\text{approximation error}}$$



Bias – Variance Tradeoff

$$\begin{aligned}
 \mathbb{E}_{X,Y} \left[(\mathbb{E}_D[\hat{f}_n(X)] - Y)^2 \right] &= \mathbb{E}_{X,Y} \left[(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X) - \epsilon)^2 \right] \\
 &= \mathbb{E}_{X,Y} \left[(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X))^2 + \epsilon^2 \right. \\
 &\quad \left. - 2\epsilon(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X)) \right] \\
 &= \mathbb{E}_{X,Y} \left[(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X))^2 \right] + \mathbb{E}_{X,Y} [\epsilon^2] \\
 &\quad - 2\mathbb{E}_{X,Y} \left[\epsilon(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X)) \right]
 \end{aligned}$$

0 since noise is independent and zero mean

$$= \underbrace{\mathbb{E}_{X,Y} \left[(\mathbb{E}_D[\hat{f}_n(X)] - f^*(X))^2 \right]}_{\text{bias}^2 - \text{how our predictor differs from the optimal}} + \underbrace{\mathbb{E}_{X,Y} [\epsilon^2]}_{\text{noise variance}}$$

Excess Risk = $\mathbb{E}_D[R(\hat{f}_n)] - R^* = \text{bias}^2 + \text{variance}$

Bottom Line

- Approximation error
 - From using a restrictive family of classifiers
- Estimation error
 - From not using best classifier in family
- Approximation and Estimation error combine to make Bias
- Variance
 - inherent to distribution

Model Selection

- Goal : minimize generalization error
- Philosophy : Occam's Razor, K.I.S.S, etc.
- Hold-out approach
 - Have training and validation data.
 - Pick model that does best on validation data
- Fancier Approaches
 - Penalize models that are likely to overfit
 - Several ways to do this

Hold-out method

Goal : pick model to minimize generalization (True) error.

Idea : train models on some data, judge them based on the rest

Hold - out procedure:

n data points available $D \equiv \{X_i, Y_i\}_{i=1}^n$

1) Split into two sets: Training dataset Validation dataset **NOT test Data !!**
 $D_T = \{X_i, Y_i\}_{i=1}^m$ $D_V = \{X_i, Y_i\}_{i=m+1}^n$

2) Use D_T for training a predictor from each model class:

$$\hat{f}_\lambda = \arg \min_{f \in \mathcal{F}_\lambda} \hat{R}_T(f)$$

 Evaluated on training dataset D_T

Hold-out method

3) Use D_v to select the model class which has smallest empirical error on D_v

$$\hat{\lambda} = \arg \min_{\lambda \in \Lambda} \hat{R}_V(f_{\lambda})$$

 Evaluated on validation dataset D_v

4) Hold-out predictor

$$\hat{f} = f_{\hat{\lambda}}$$

Intuition: Small error on one set of data will not imply small error on a randomly sub-sampled second set of data

Ensures method is “stable”

Hold-out method

Drawbacks:

- May not have enough data to afford setting one subset aside for getting a sense of generalization abilities
- Validation error may be misleading (bad estimate of generalization error) if we get an “unfortunate” split

Limitations of hold-out can be overcome by a family of random sub-sampling methods at the expense of more computation.

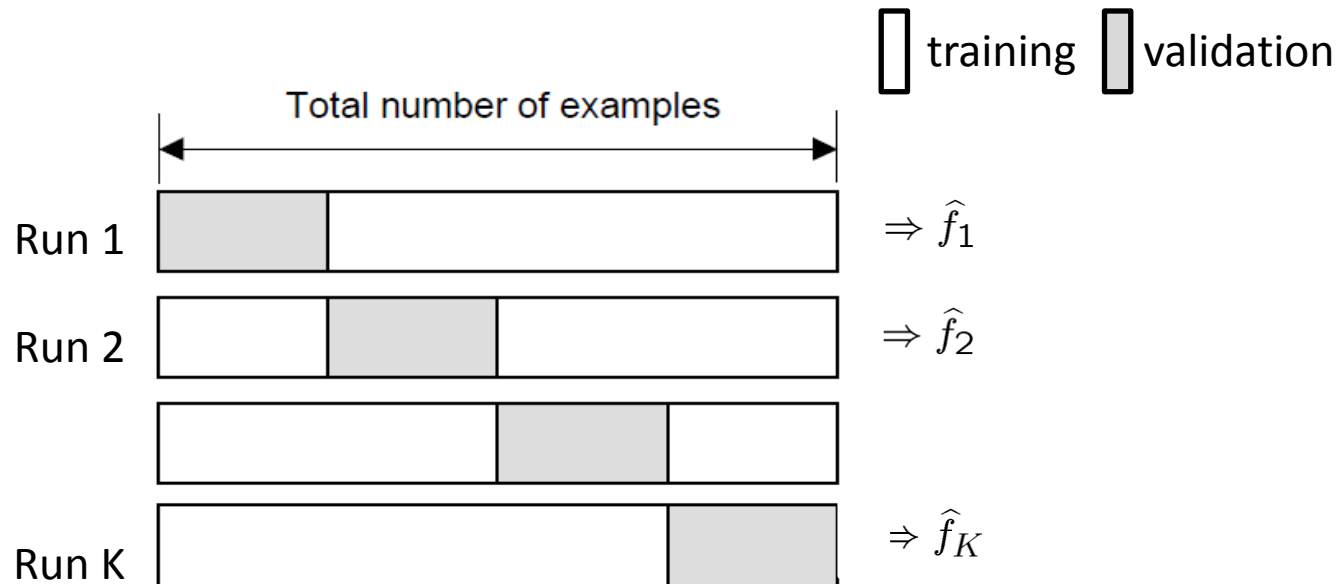
Cross-validation

K-fold cross-validation

Create K-fold partition of the dataset.

Form K hold-out predictors, each time using one partition as validation and rest K-1 as training datasets.

Final predictor is average/majority vote over the K hold-out estimates.




For each model family, train K models and average their errors
Pick the model family that performed best
Re-train this family on all of the data

Fancier Approaches

- Minimizing empirical risk is, um, risky
- Add a “penalty term” of some sort to the empirical risk
- This term should indicate how prone a model is to overfitting

Structural Risk Minimization

Recall : Best model minimizes true risk



Problem : We don't know true error


Idea : Bound the true error as a function of complexity and minimize that

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \left\{ \hat{R}_n(f) + C(f) \right\}$$

└─→ Bound on deviation from true risk

With high probability, $|R(f) - \hat{R}_n(f)| \leq C(f) \quad \forall f \in \mathcal{F}$ Concentration bounds (later)

Structural Risk Minimization



In practice, theoretical bounds are way too high, so we scale them

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \{ \hat{R}_n(f) + \lambda C(f) \}$$

Choose by cross-validation!

Complexity Regularization

Penalize complex models using **prior knowledge**.

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \left\{ \hat{R}_n(f) + C(f) \right\}$$

Cost of model
(log prior)

Bayesian viewpoint:

prior probability of $f \equiv e^{-C(f)}$

cost is small if f is highly probable, cost is large if f is improbable

ERM (empirical risk minimization) over a restricted class F , e.g. linear classifiers,
 \equiv uniform prior on $f \in F$, zero probability for other predictors

$$\hat{f}_n^L = \arg \min_{f \in \mathcal{F}_L} \hat{R}_n(f)$$

Note the common pattern

- Start with MLE
 - Minimize $-(\text{probability of data})$
 - But this overfits
- So we add a penalty term for “bad-looking” models
 - Least squares
 - Model complexity
- This has a Bayesian interpretation

Information Criteria

Penalize complex models based on their **information content**.

$$\hat{f}_n = \arg \min_{f \in \mathcal{F}} \left\{ \hat{R}_n(f) + C(f) \right\}$$

MDL (Minimum Description Length)

\rightarrow # bits needed to describe f
(description length)

Example: Binary Decision trees $\mathcal{F}_k^T = \{\text{tree classifiers with } k \text{ leafs}\}$


$\mathcal{F}^T = \bigcup_{k \geq 1} \mathcal{F}_k^T$ prefix encode each element f of \mathcal{F}^T

$$C(f) = 3k - 1 \text{ bits}$$

k leafs $\Rightarrow 2k - 1$ nodes

$2k - 1$ bits to encode tree structure

+ k bits to encode label of each leaf (0/1)



5 leafs \Rightarrow 9 bits to encode structure